

# Specialist Programme on Artificial Intelligence for IT & ITES Industry

## *Convolutional neural networks, Part 2*

By Dr. Tan Jen Hong  
[tanjenhong@nus.edu.sg](mailto:tanjenhong@nus.edu.sg)

Singapore e-Government Leadership Centre  
National University of Singapore

© 2020 NUS. The contents contained in this document may not be reproduced in any form or by any means, without the written permission of ISS, NUS other than for the purpose for which it has been supplied.

Inspire


Lead

Transform

**The need of a deeper network and  
functional APIs**

# Representation

Do you think deep neural net can represent this function?




Source: By Brendan Fortuner

# Universal approximation theorem

A theorem for all problem?

- A feedforward network with a single layer is enough to represent any function
- What is the implication?
- With the theorem, is deep neural network the panacea for all our challenges?



Source: By Brendan Fortuner

# Back to the real world

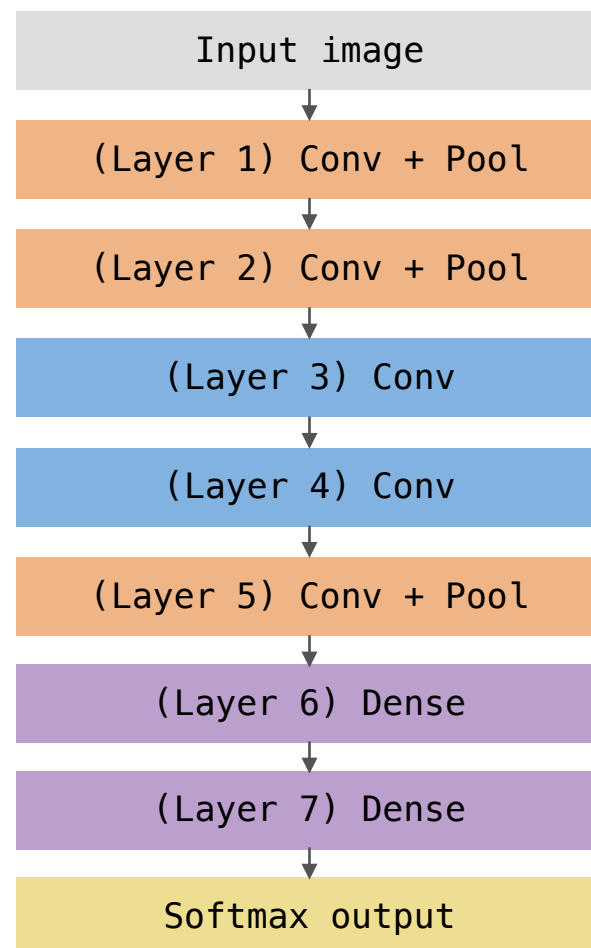
The issue at hand

- Although by universal approximation theorem, possible to perform good classification on a single layer net, so far no learning algorithm can achieve that
- Key: a single layer net can approximate any function (proven), but the theorem does not provide any clue to achieve that
- On the contrary, the experiences from the past decade show that depth of a net is the key to great performance

# The importance of depth

A study on Krizhevsky et al. model (2012)

- The model that won ILSVRC 2012
- 8 layers in total, 60 million parameters, 650,000 neurons
- Re-implementation gives 18.1% top-5 error



Source: Re-implementation by Rob Fergus

# Top-5 error

How about top-1 error



Persian cat

Source: [http://www.vetstreet.com/cats/persian#1\\_ugw20zmq](http://www.vetstreet.com/cats/persian#1_ugw20zmq)

## Top-1 error

Chihuahua (0.4)  
Hyena (0.25)  
Koala (0.15)  
Persian cat (0.1)  
Burmese cat (0.02)

Considered **incorrect**

## Top-5 error

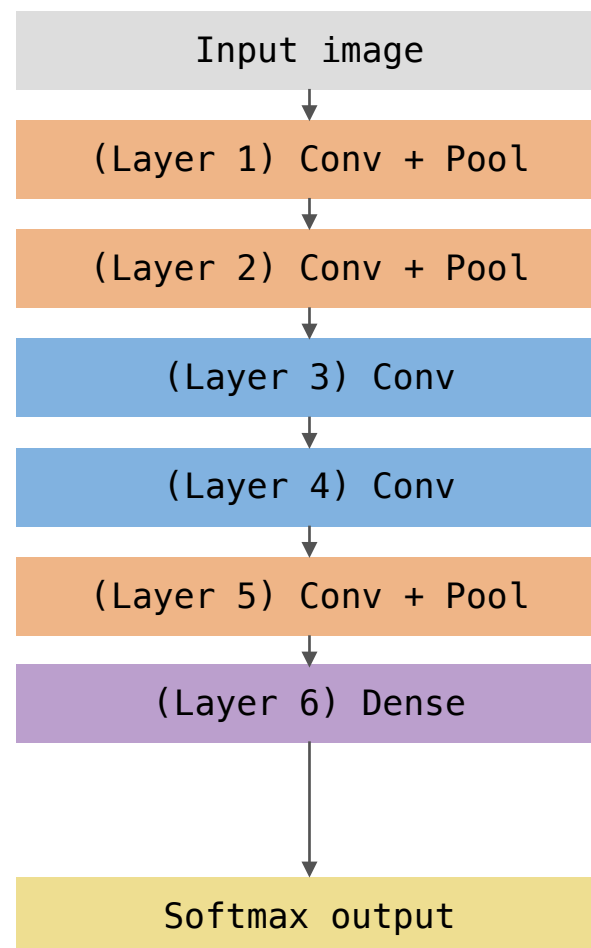
Chihuahua (0.4)  
Hyena (0.25)  
Koala (0.15)  
Persian cat (0.1)  
Burmese cat (0.02)

Considered **correct**

# The importance of depth

A study on Krizhevsky et al. model (2012)

- Remove layer 7, the fully connected layer
- 16 million less parameters compared to the original model
- Only 1.1% drop in performance!



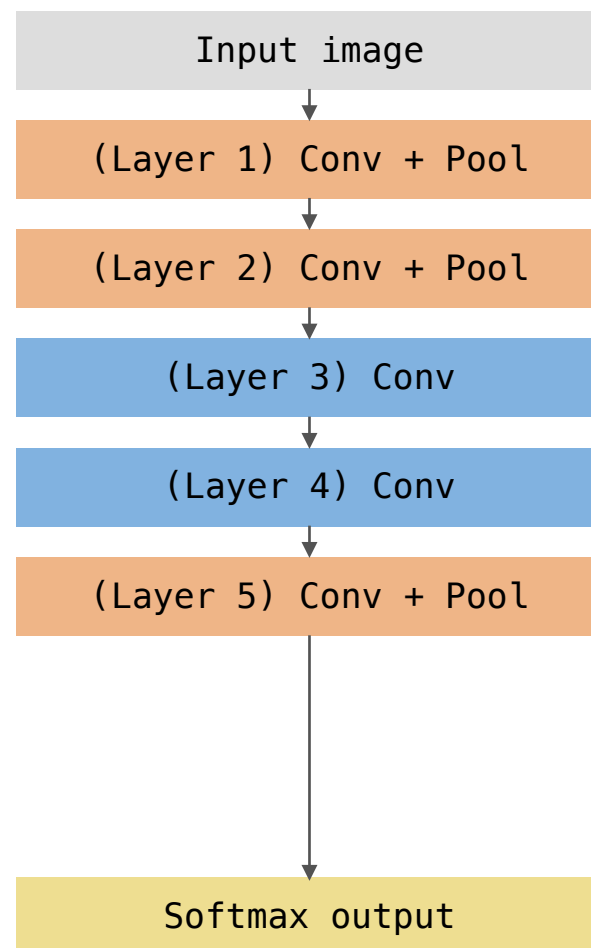
Source: Re-implementation by Rob Fergus



# The importance of depth

A study on Krizhevsky et al. model (2012)

- Remove both layer 6 and layer 7, the two fully connected layers
- 50 million less parameters compared to the original model
- 5.7% drop in performance




Source: Re-implementation by Rob Fergus

# The importance of depth

A study on Krizhevsky et al. model (2012)


- Now try remove upper layers, the feature extractor. Remove layer 3 and 4
- 1 million less parameters compared to the original model
- 3.0% drop in performance



Source: Re-implementation by Rob Fergus

# The importance of depth

A study on Krizhevsky et al. model (2012)



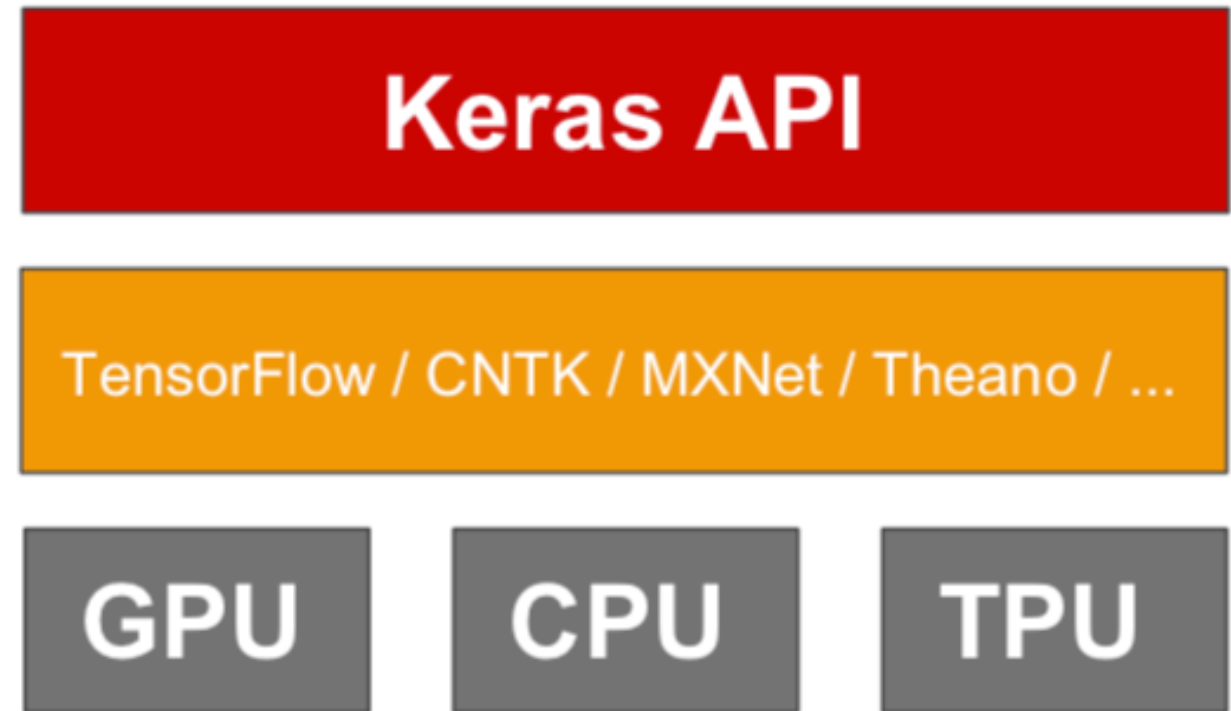
- Now try remove upper layers, the feature extractor. Remove layer 3, 4, 6, and 7
- Only 4 layers left
- **33.5%** drop in performance!
- Depth is the key

Source: Re-implementation by Rob Fergus

**Keras**

# What is Keras

The basic architecture




Source: Francois Chollet

# Keras + Tensorflow

The official high-level API of Tensorflow

- Possible to build deep net solely using Tensorflow, but toooo many repetition
- Need to define every detail (weight, bias, initialization and etc.)
- Since Tensorflow v1.4, Keras is part of the core
- In Tensorflow v1.13, we have tensorflow.keras module
- From Tensorflow v2.0, Keras API is **THE preferred way** of building neural network



Source: Francois Chollet

# Keras

## Three API styles

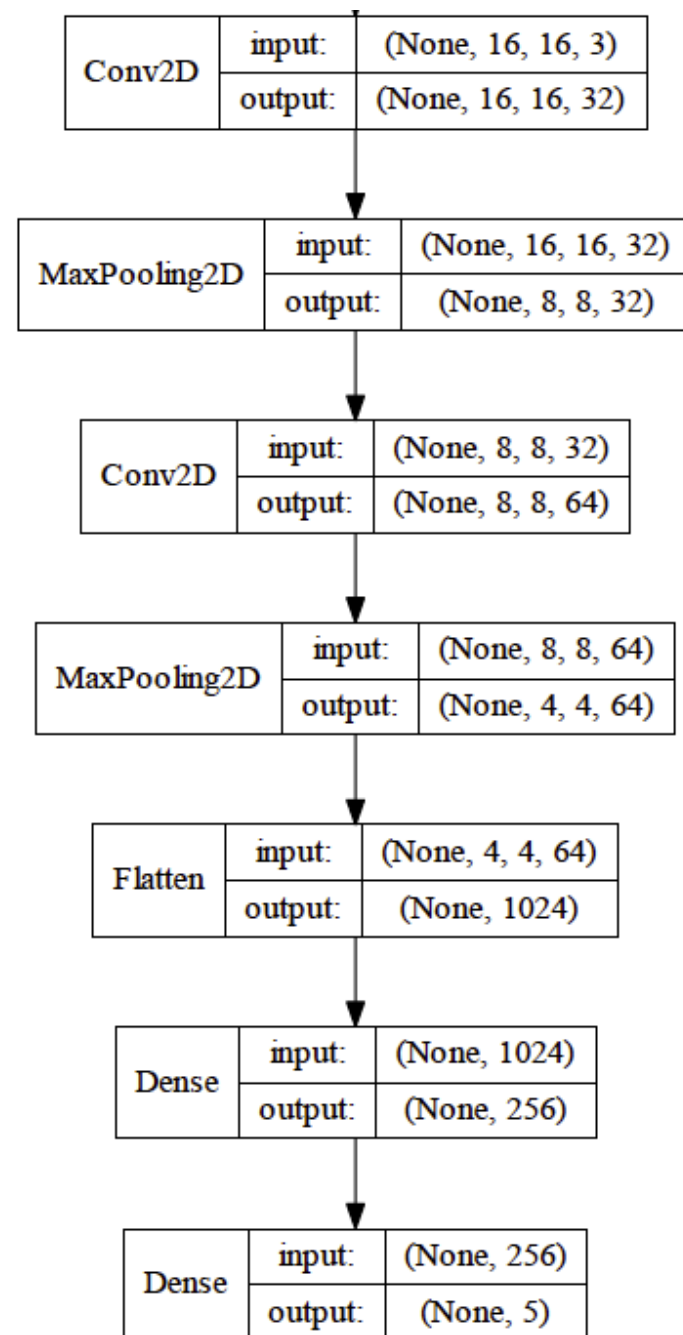
- **Sequential model**  
Super simple  
Only for single-input, single-output  
sequential layer stacks  
Good for 70+% of use cases
- **Functional API**
  - Works like playing Lego bricks
  - Multi-input, multi-output, arbitrary static graph topologies
  - Good for 95% of use cases
- **Model subclassing**  
Maximum flexibility  
Larger potential error surface

Source: Francois Chollet

# Keras

## Sequential model

- Adding layers by sequence (layer by layer)
- `input_shape` must be present in the first layer

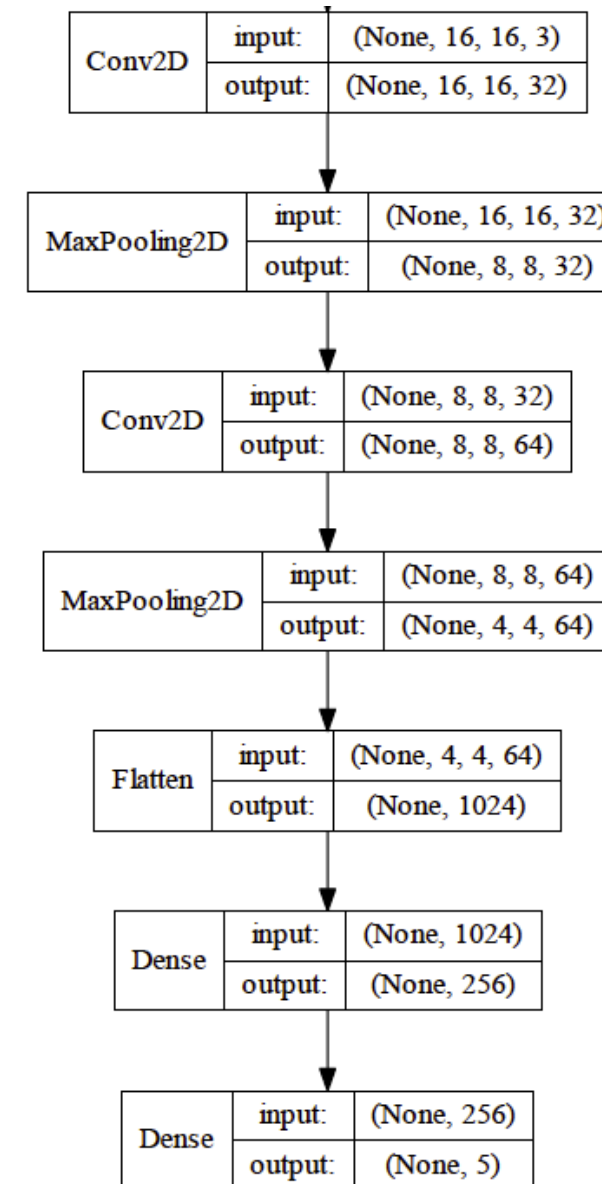




# Keras

## Sequential model

```
> def createSeqModel():  
    model = Sequential()  
    model.add(Conv2D(32, (3, 3),  
                    input_shape=(16, 16, 3),  
                    padding='same',  
                    activation='relu'))  
    model.add(MaxPooling2D(pool_size=(2, 2)))  
    model.add(Conv2D(64, (3, 3),  
                    padding='same',  
                    activation='relu'))  
    model.add(MaxPooling2D(pool_size=(2, 2)))  
  
    model.add(Flatten())  
    model.add(Dense(256, activation='relu'))  
    model.add(Dense(5, activation='softmax'))  
  
    model.compile(loss='categorical_crossentropy',  
                  optimizer='rmsprop',  
                  metrics=['accuracy'])  
  
    return model
```



# Keras

## Sequential model

- Create the model and show the model summary

```
> modelSeq = createSeqModel()  
> modelSeq.summary()
```

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 16, 16, 32)	896
max_pooling2d (MaxPooling2D)	(None, 8, 8, 32)	0
conv2d_1 (Conv2D)	(None, 8, 8, 64)	18496
max_pooling2d_1 (MaxPooling2D)	(None, 4, 4, 64)	0
flatten (Flatten)	(None, 1024)	0
dense (Dense)	(None, 256)	262400
dense_1 (Dense)	(None, 5)	1285
Total params: 283,077		
Trainable params: 283,077		
Non-trainable params: 0		

# Keras

## Sequential model

- The problem

Not possible for multiple input

Not possible for multiple output

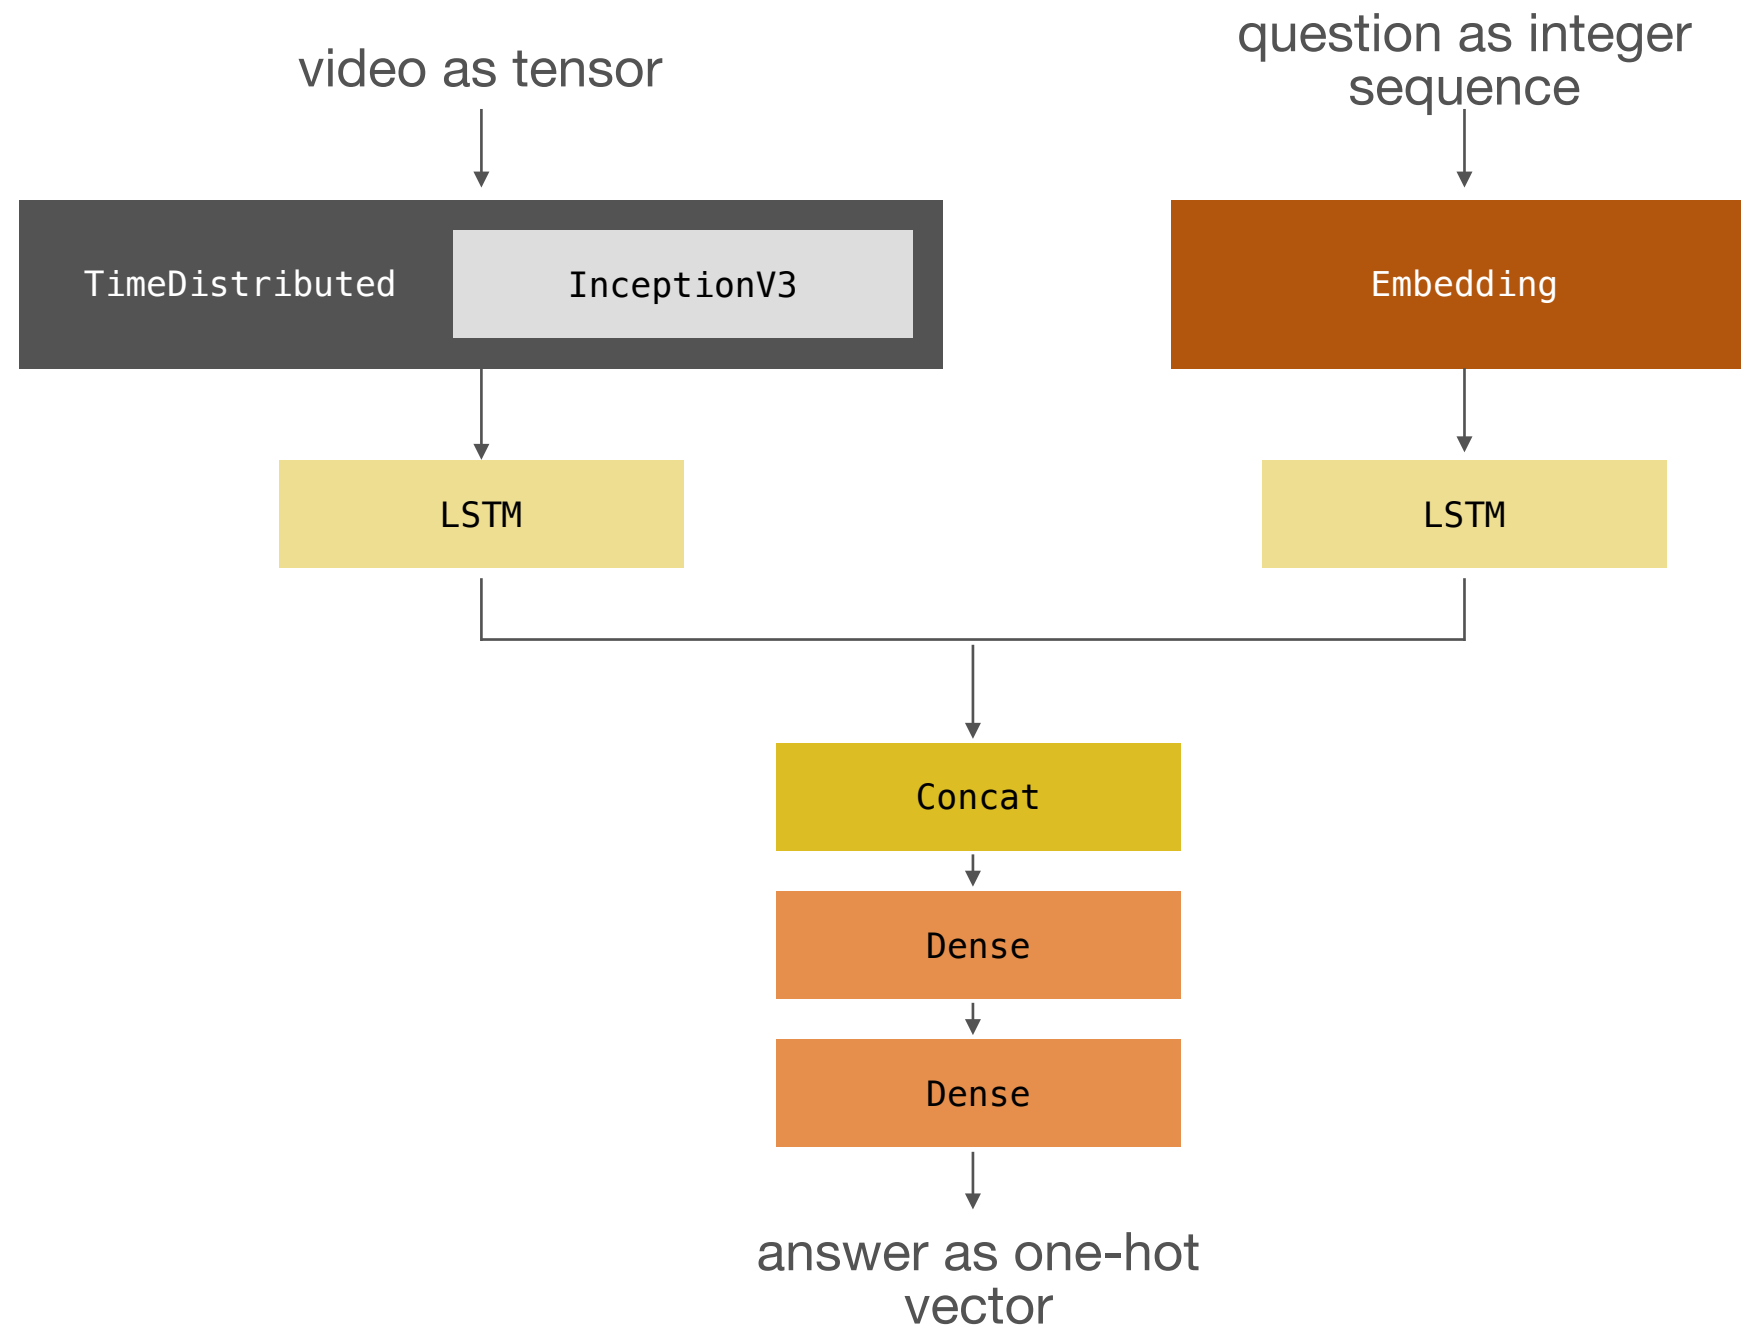
A single direction of flow of tensors,  
no branching, no merging

Not possible to reuse layer; shared  
layer is not possible

```
> def createSeqModel():  
    model = Sequential()  
    model.add(Conv2D(32, (3, 3),  
                    input_shape=(16, 16, 3),  
                    padding='same',  
                    activation='relu'))  
    model.add(MaxPooling2D(pool_size=(2, 2)))  
    model.add(Conv2D(64, (3, 3),  
                    padding='same',  
                    activation='relu'))  
    model.add(MaxPooling2D(pool_size=(2, 2)))  
  
    model.add(Flatten())  
    model.add(Dense(256, activation='relu'))  
    model.add(Dense(5, activation='softmax'))  
  
    model.compile(loss='categorical_crossentropy',  
                  optimizer='rmsprop',  
                  metrics=['accuracy'])  
    return model
```


# Multiple inputs

An video and a question



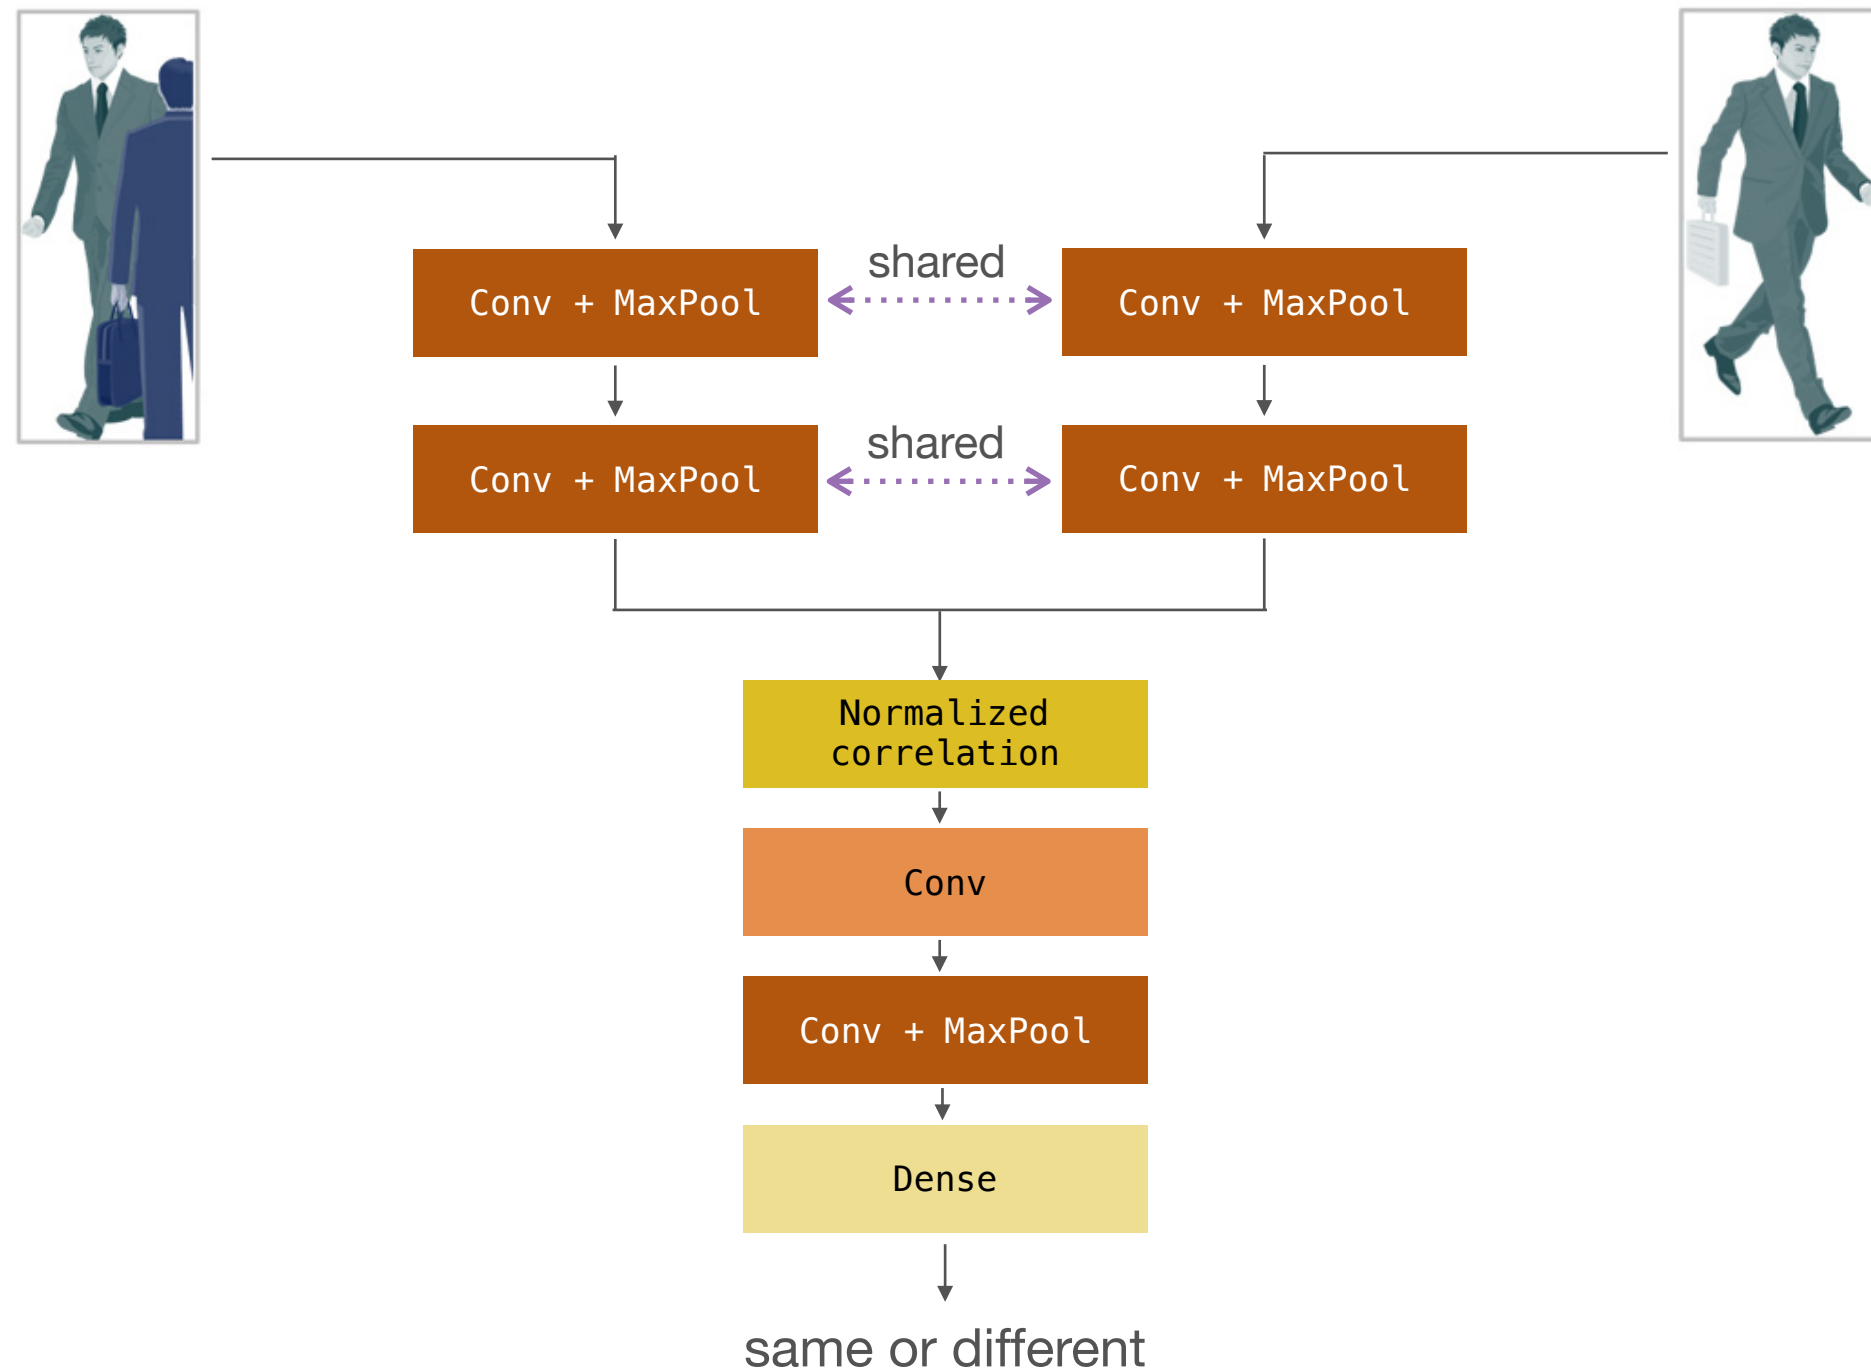
# Multiple outputs

An adjusted image and a segmentation



# Shared layers

## Person re-identification



# Keras

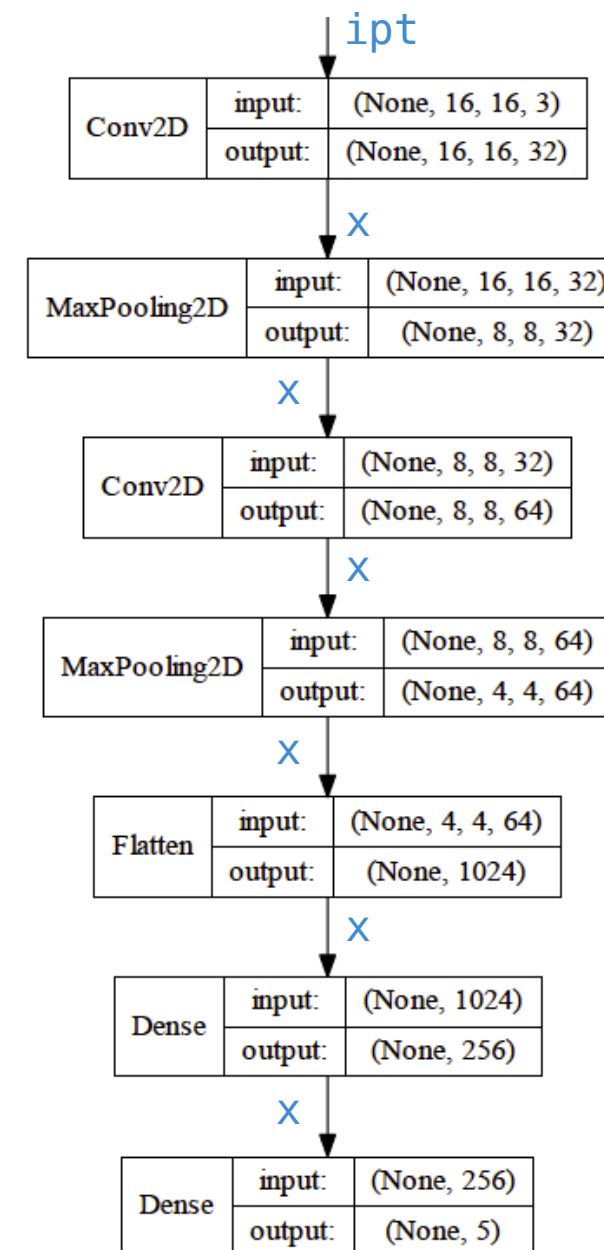
## functional APIs

```
> def createFuncModel():
    ipt      = Input(shape=(16,16,3))
    x        = Conv2D(32,(3,3),
                      padding='same',
                      activation='relu')(ipt)
    x        = MaxPooling2D(pool_size=(2,2))(x)
    x        = Conv2D(64,(3,3),
                      padding='same',
                      activation='relu')(x)
    x        = MaxPooling2D(pool_size=(2,2))(x)
    x        = Flatten()(x)
    x        = Dense(256,activation='relu')(x)
    x        = Dense(5,activation='relu')(x)

    model    = Model(inputs=ipt,outputs=x)

    model.compile(loss='categorical_crossentropy',
                  optimizer='rmsprop',
                  metrics=['accuracy'])

    return model
```



# Keras

## Comparison

### functional APIs

```
> def createFuncModel():
    ipt      = Input(shape=(16,16,3))
    x        = Conv2D(32,(3,3),
                      padding='same',
                      activation='relu')(ipt)
    x        = MaxPooling2D(pool_size=(2,2))(x)
    x        = Conv2D(64,(3,3),
                      padding='same',
                      activation='relu')(x)
    x        = MaxPooling2D(pool_size=(2,2))(x)
    x        = Flatten()(x)
    x        = Dense(256,activation='relu')(x)
    x        = Dense(5,activation='relu')(x)

    model    = Model(inputs=ipt,outputs=x)

    model.compile(loss='categorical_crossentropy',
                  optimizer='rmsprop',
                  metrics=['accuracy'])

    return model
```

### Sequential model

```
> def createSeqModel():
    model    = Sequential()
    model.add(Conv2D(32,(3,3),
                     input_shape=(16,16,3),
                     padding='same',
                     activation='relu'))
    model.add(MaxPooling2D(pool_size=(2,2)))
    model.add(Conv2D(64,(3,3),
                     padding='same',
                     activation='relu'))
    model.add(MaxPooling2D(pool_size=(2,2)))

    model.add(Flatten())
    model.add(Dense(256,activation='relu'))
    model.add(Dense(5,activation='softmax'))

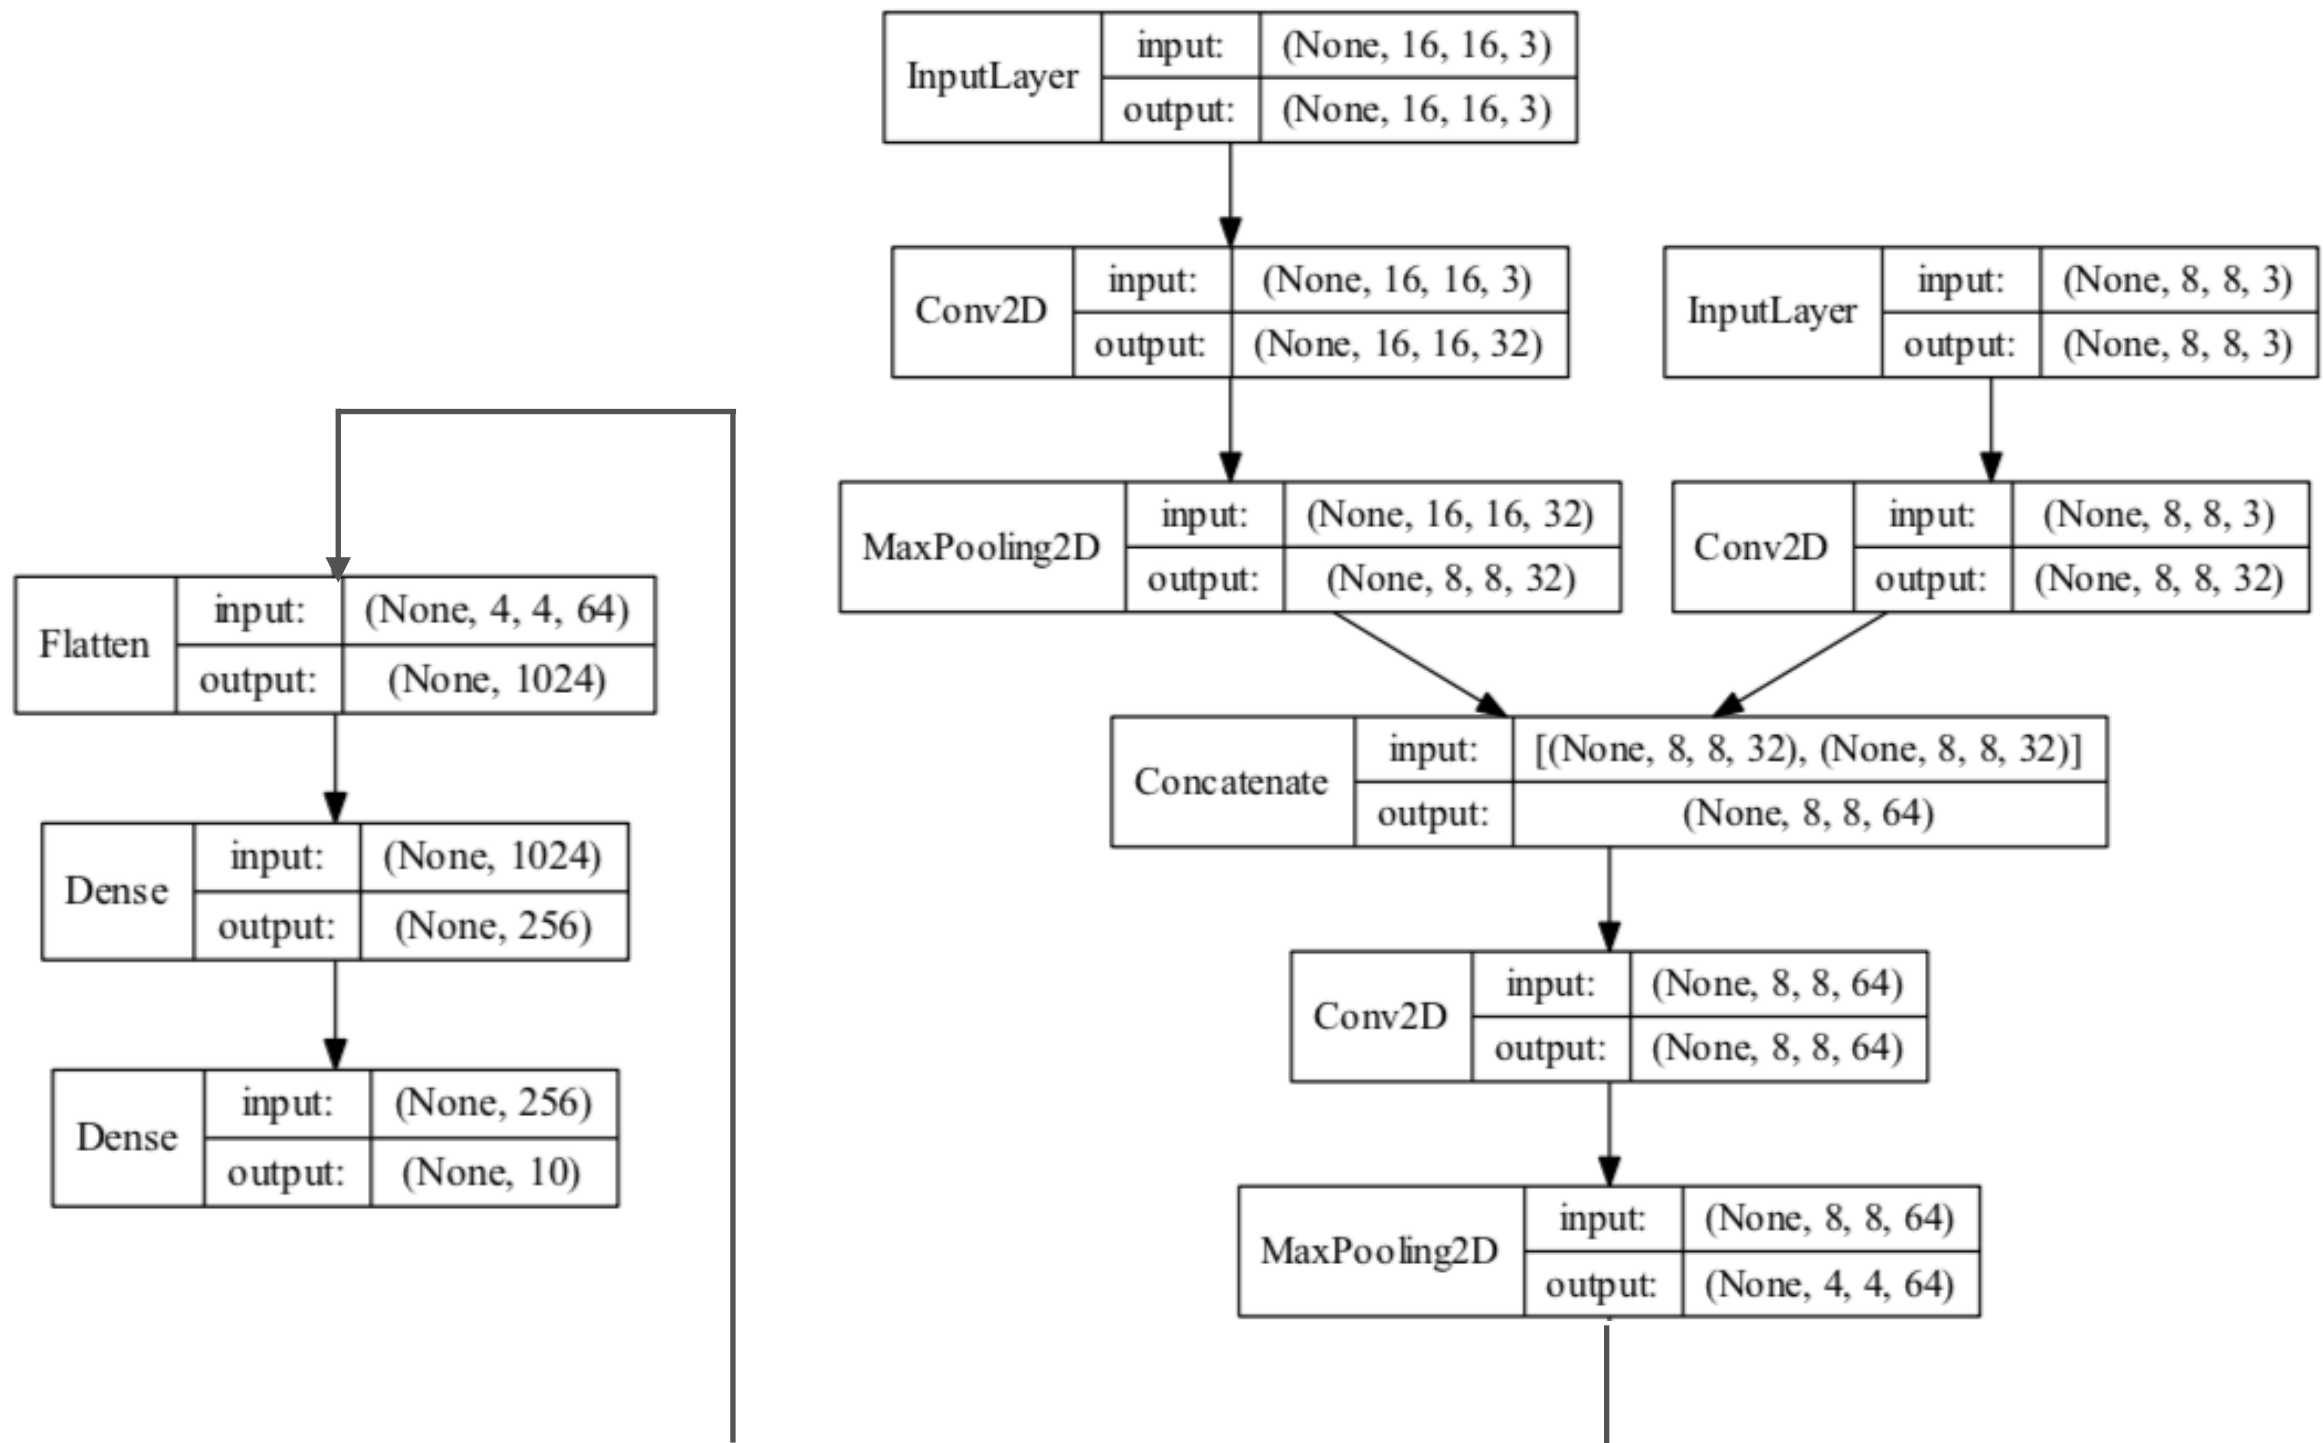
    model.compile(loss='categorical_crossentropy',
                  optimizer='rmsprop',
                  metrics=['accuracy'])

    return model
```



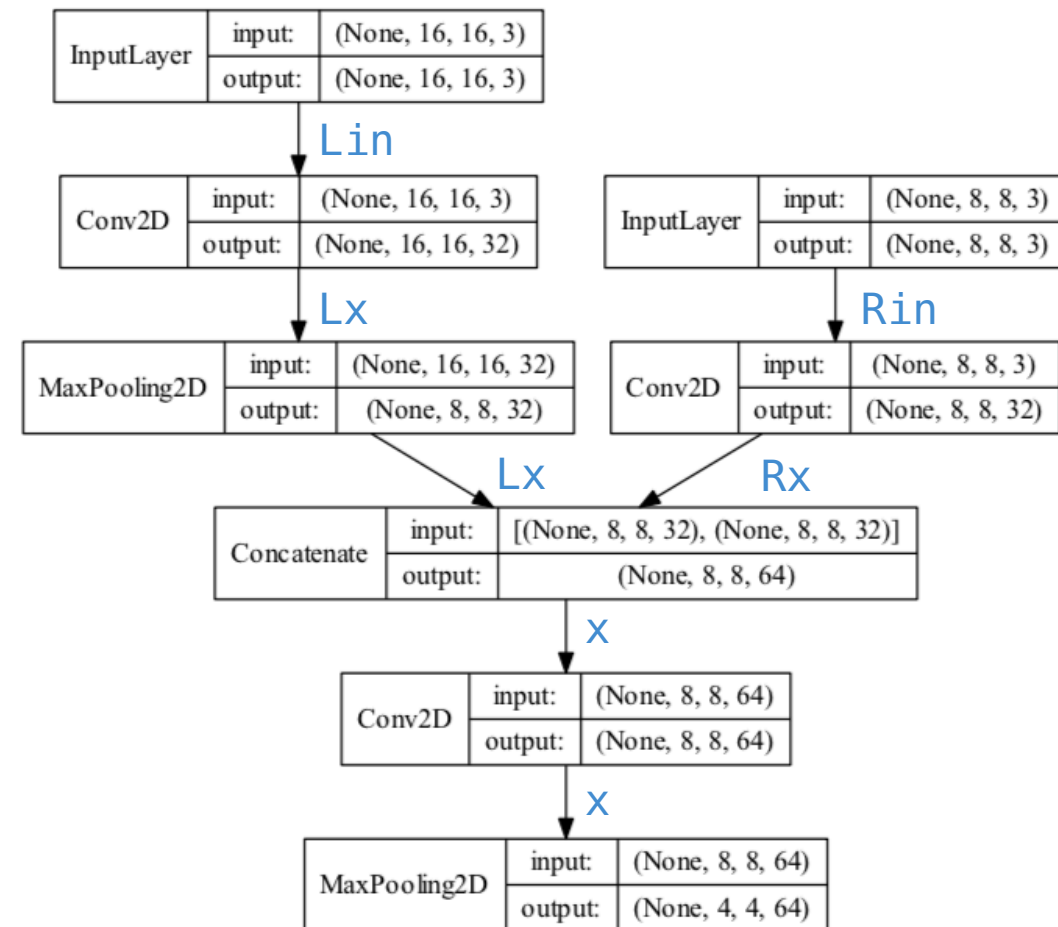
# 'Y' shape architecture

how to create?



# 'Y' shape architecture

the code



```

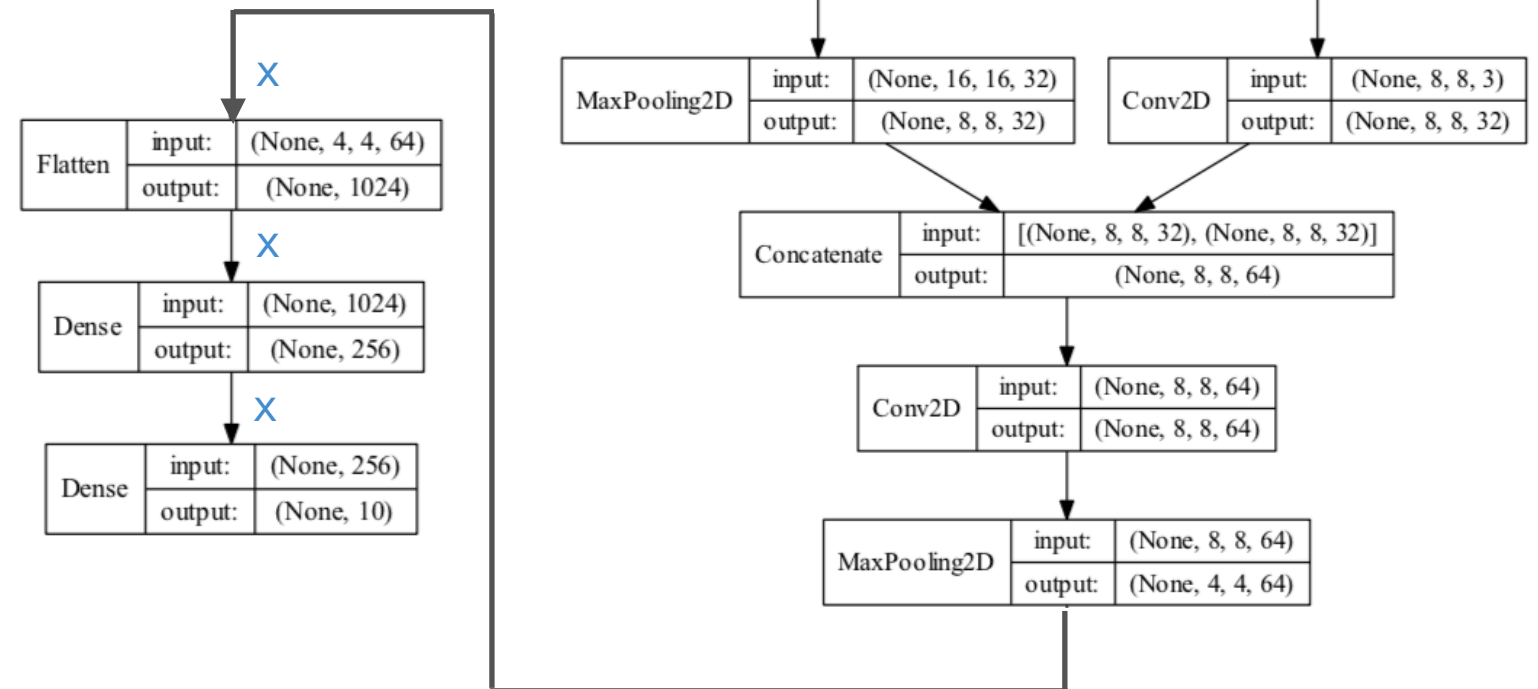
> def createDualInputModel():
    Lin    = Input(shape=(16,16,3))
    Lx     = Conv2D(32,(3,3),padding='same',activation='relu')(Lin)
    Lx     = MaxPooling2D(pool_size=(2,2))(Lx)

    Rin    = Input(shape=(8,8,3))
    Rx     = Conv2D(32,(3,3),padding='same',activation='relu')(Rin)

    x      = concatenate([Lx,Rx],axis=-1)
    x      = Conv2D(64,(3,3),padding='same',activation='relu')(x)
    x      = MaxPooling2D(pool_size=(2,2))(x)
    ...
  
```

# 'Y' shape architecture

the code



```
> def createDualInputModel():
    ....
    x      = Flatten()(x)
    x      = Dense(256,activation='relu')(x)
    x      = Dense(10,activation='softmax')(x)

    model  = Model(inputs=[Lin,Rin],outputs=x)
    model.compile(loss='categorical_crossentropy',
                  optimizer='rmsprop',
                  metrics=['accuracy'])

    return model
```

# 'Y' shape architecture

Putting all together ....

- Remember to include the below before the function

```
from tensorflow.keras.layers import concatenate
```

```
> def createDualInputModel():
    Lin      = Input(shape=(16,16,3))
    Lx       = Conv2D(32,(3,3),padding='same',activation='relu')(Lin)
    Lx       = MaxPooling2D(pool_size=(2,2))(Lx)

    Rin      = Input(shape=(8,8,3))
    Rx       = Conv2D(32,(3,3),padding='same',activation='relu')(Rin)

    x        = concatenate([Lx,Rx],axis=-1)
    x        = Conv2D(64,(3,3),padding='same',activation='relu')(x)
    x        = MaxPooling2D(pool_size=(2,2))(x)

    x        = Flatten()(x)
    x        = Dense(256,activation='relu')(x)
    x        = Dense(10,activation='softmax')(x)

    model    = Model(inputs=[Lin,Rin],outputs=x)
    model.compile(loss='categorical_crossentropy',
                  optimizer='rmsprop',
                  metrics=['accuracy'])

    return model
```

# Multiple inputs

the code

- Create the model and show the model summary

```
> modelDual = createDualInputModel()
> modelDual.summary()
```

Layer (type)	Output Shape	Param #	Connected to
=====			
input_1 (InputLayer)	(None, 16, 16, 3)	0	
conv2d (Conv2D)	(None, 16, 16, 32)	896	input_1[0][0]
input_2 (InputLayer)	(None, 8, 8, 3)	0	
max_pooling2d (MaxPooling2D)	(None, 8, 8, 32)	0	conv2d[0][0]
conv2d_1 (Conv2D)	(None, 8, 8, 32)	896	input_2[0][0]
concatenate (Concatenate)	(None, 8, 8, 64)	0	max_pooling2d[0][0] conv2d_1[0][0]
conv2d_2 (Conv2D)	(None, 8, 8, 64)	36928	concatenate[0][0]
max_pooling2d_1 (MaxPooling2D)	(None, 4, 4, 64)	0	conv2d_2[0][0]
flatten (Flatten)	(None, 1024)	0	max_pooling2d_1[0][0]
dense (Dense)	(None, 128)	131200	flatten[0][0]
dense_1 (Dense)	(None, 3)	387	dense[0][0]
=====			
Total params: 170,307			
Trainable params: 170,307			
Non-trainable params: 0			

# Multiple inputs

training

- Assume RDat and LDat are the training input, TLbl is the label for the training
- Furthermore, vRDat and vLDat is the validation input, and vLbl is the label
- The training is done by

```
> model.fit([RDat,LDat],  
            TLbl,  
            validation_data=( [vRDat,vLDat], vLbl),  
            epochs=100,  
            batch_size=128,  
            shuffle=True,  
            callbacks=callbacks_list)
```

# Multiple inputs

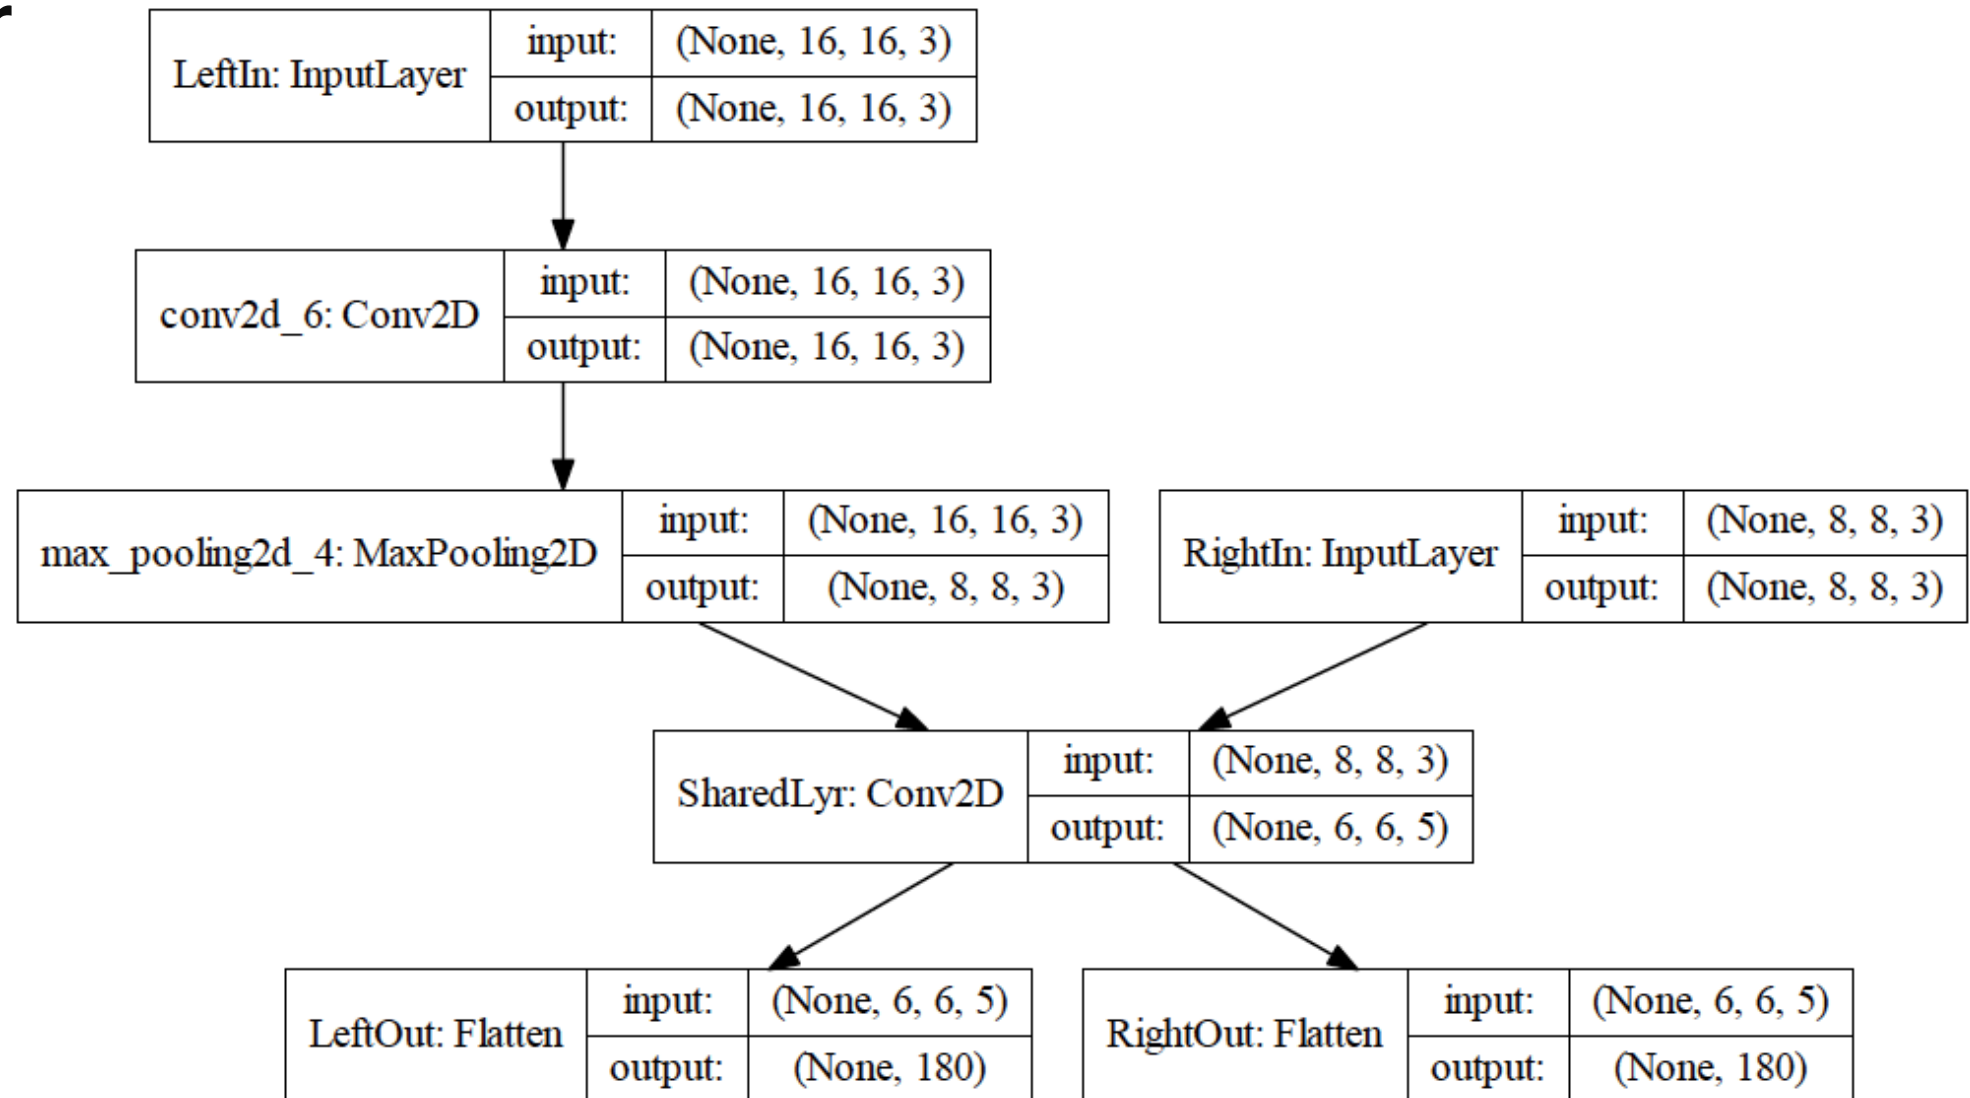
## Getting the model plot

- To plot model, it requires pydot and graphviz
- On Mac, there is a need for additional installation for graphviz

```
> from tensorflow.keras.utils import plot_model  
  
> plot_model(modelDual,  
             to_file='Dual_model.pdf',  
             show_shapes=True,  
             show_layer_names=False,  
             rankdir='TB')
```

# Shared layer

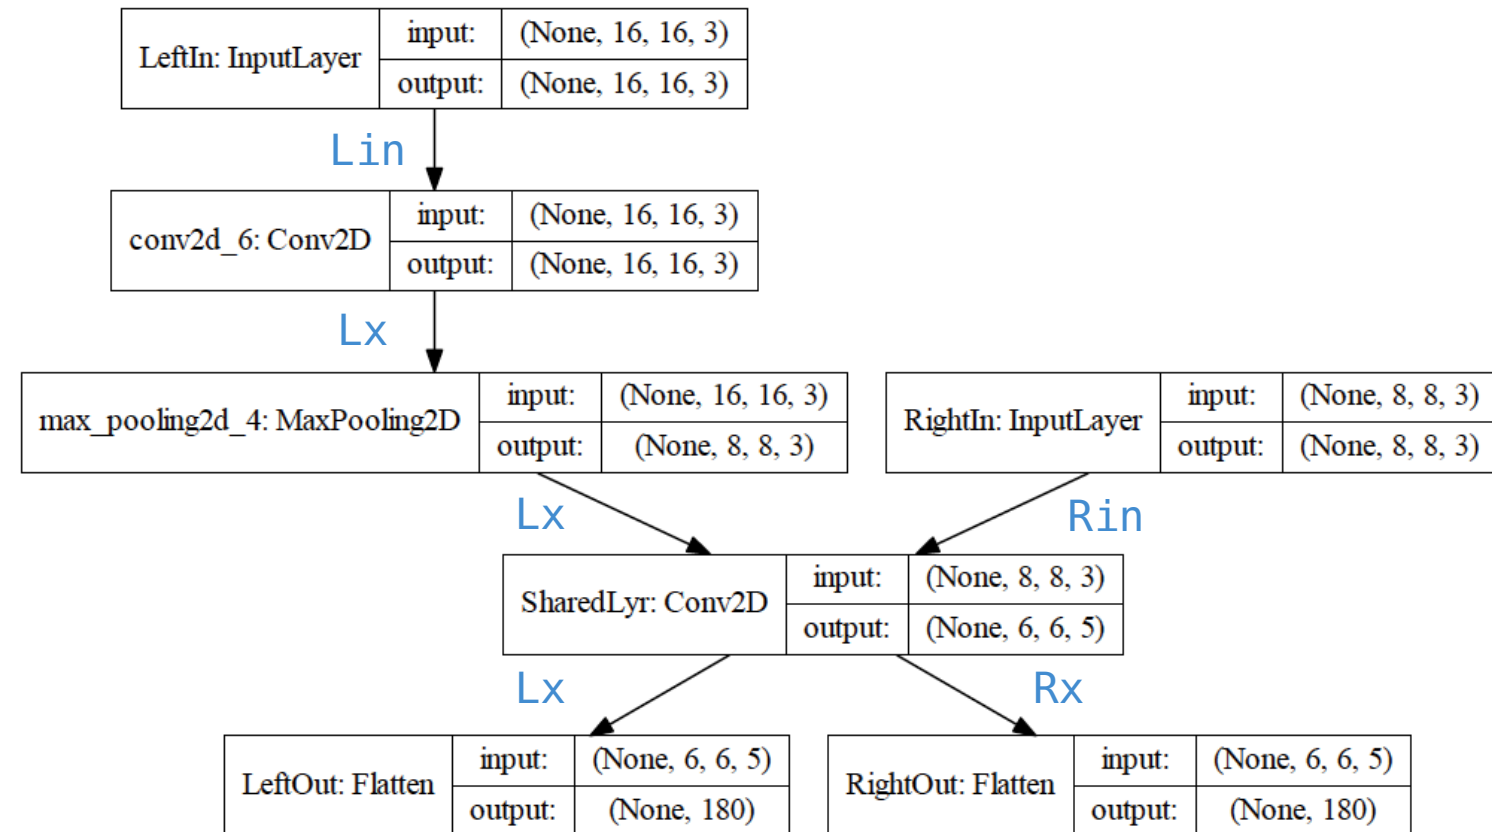
how to create?





# Shared layer

the code



```

> def createSharedModel():
    shared = Conv2D(5, (3,3), activation='relu', name='SharedLyr')

    Lin = Input(shape=(16,16,3), name='LeftIn')
    Lx = Conv2D(3, (3,3), padding='same', activation='relu')(Lin)
    Lx = MaxPooling2D(pool_size=(2,2))(Lx)
    Lx = shared(Lx)
    Lx = Flatten(name='LeftOut')(Lx)

    Rin = Input(shape=(8,8,3), name='RightIn')
    Rx = shared(Rin)
    Rx = Flatten(name='RightOut')(Rx)
    model = Model(inputs=[Lin,Rin], outputs=[Lx,Rx])

    ...
  
```

# Shared layer

the code

- For model with more than one output, pass in a dictionary that specifies the loss function for each output

```
> def createSharedModel():
    shared = Conv2D(5, (3,3), activation='relu', name='SharedLyr')

    Lin = Input(shape=(16,16,3), name='LeftIn')
    Lx = Conv2D(3, (3,3), padding='same', activation='relu')(Lin)
    Lx = MaxPooling2D(pool_size=(2,2))(Lx)
    Lx = shared(Lx)
    Lx = Flatten(name='LeftOut')(Lx)

    Rin = Input(shape=(8,8,3), name='RightIn')
    Rx = shared(Rin)
    Rx = Flatten(name='RightOut')(Rx)

    model = Model(inputs=[Lin,Rin], outputs=[Lx,Rx])
    model.compile(loss={'LeftOut': 'categorical_crossentropy',
                        'RightOut': 'mean_squared_error'},
                  optimizer='rmsprop',
                  metrics=['accuracy'])

    return model
```

# Keras

## Merge layer

- Other possible layers to fuse multiple flows of tensors in Keras

```
tensorflow.keras.layers.add  
tensorflow.keras.layers.subtract  
tensorflow.keras.layers.multiply  
tensorflow.keras.layers.average  
tensorflow.keras.layers.maximum  
tensorflow.keras.layers.minimum
```

- For the above, the tensor inputs to the layer should have the same size
- Good enough for most use cases, otherwise write your own lambda layer to perform the necessary task

# Keras

## Multiple GPU

- Training using multiple GPUs is easy in Keras
- In the example, each GPU will process 64 samples in a single batch
- Why a need for multiple GPUs?

```
> from tensorflow.keras.utils import multi_gpu_model  
  
> parallel = multi_gpu_model(model, gpus=4)  
> parallel.compile(loss='categorical_crossentropy',  
                  optimizer='rmsprop')  
  
> parallel.fit(x, y, epochs=50, batch_size=256)
```

**Singapore e-Government Leadership Centre  
National University of Singapore  
29 Heng Mui Keng Terrace  
Block D & E  
Singapore 119620**

**Tel : (65) 6516 1156  
Fax : (65) 6778 2571  
URL : [www.egl.sg](http://www.egl.sg)  
Email : [egl-enquiries@nus.edu.sg](mailto:egl-enquiries@nus.edu.sg)**

*Inspire* *Lead* *Transform*